Jupyter Architecture

An overview

What is going on under the hood?
Let's start with the Browser.

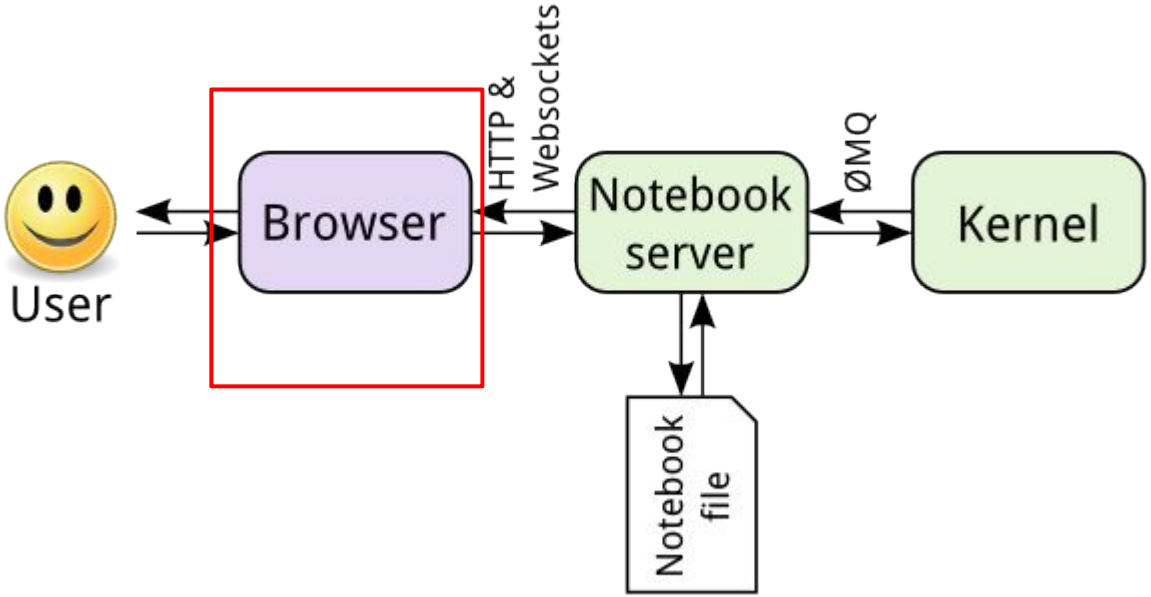

Browser

- You are viewing/editing a document


```
{
 "cells": [...],
 "metadata": {
  "kernelspec": { "display_name": "Python 3", "language": "python", "name": "python3"},

   ...
  }
 },
 "nbformat": 4,
 "nbformat_minor": 2
}
```

Browser

- Jupyter’s Notebook frontend handles the following
  - User input
  - Rendering (Markdown, HTML, Javascript)
  - Sending changes to the Notebook Server
  - Sending code execute requests to the Notebook Server

Browser

- Communicates using HTTP and Websockets
  - Websockets allow for a 2-way link between your browser and the server
  - Network client

Notebook Server

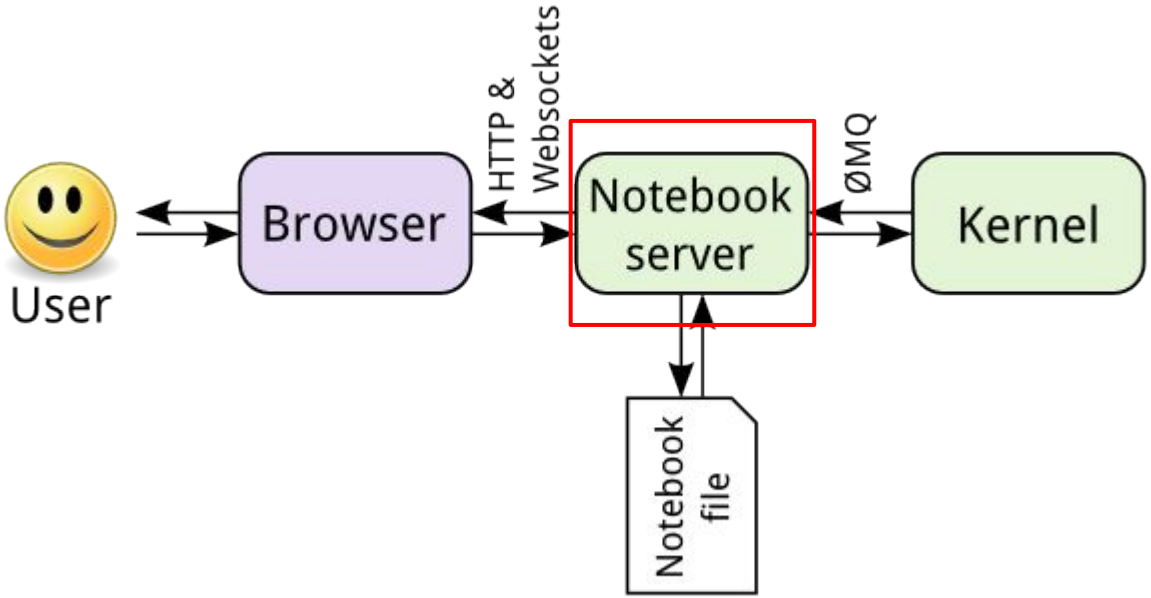

Notebook Server

- Responsibilities
  - Saves/Loads notebook files
  - Spawns a kernel for you
  - Shuts down kernels
  - Communicates with the kernel
    - uses the ZeroMQ (Message Queue) protocol

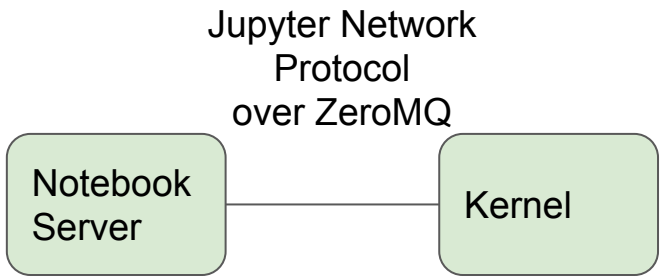

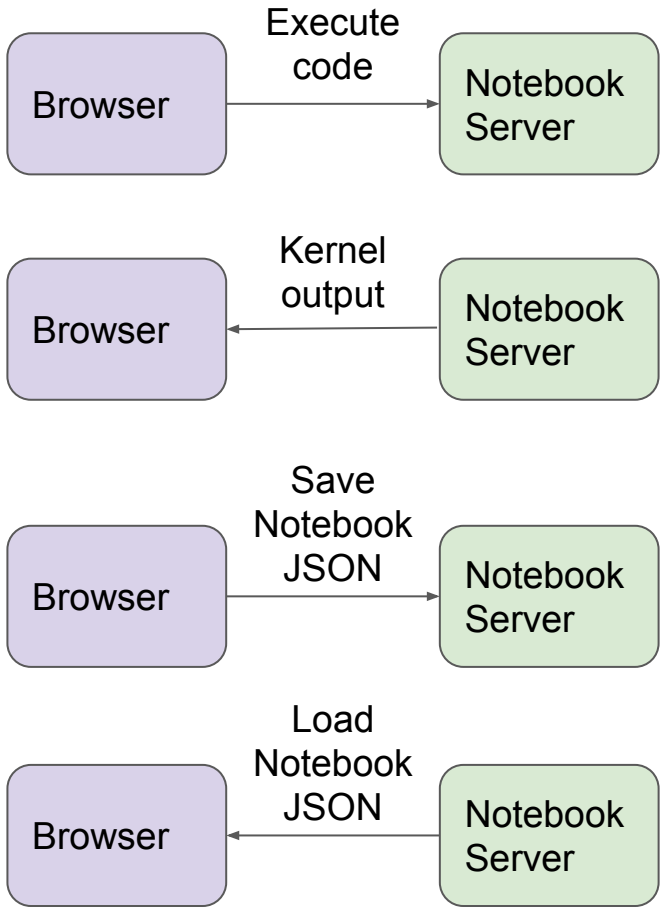

Notebook Server

- REST API
- Handles Notebook File Menu requests (Download as)
- IPyparallel cluster management

Kernel

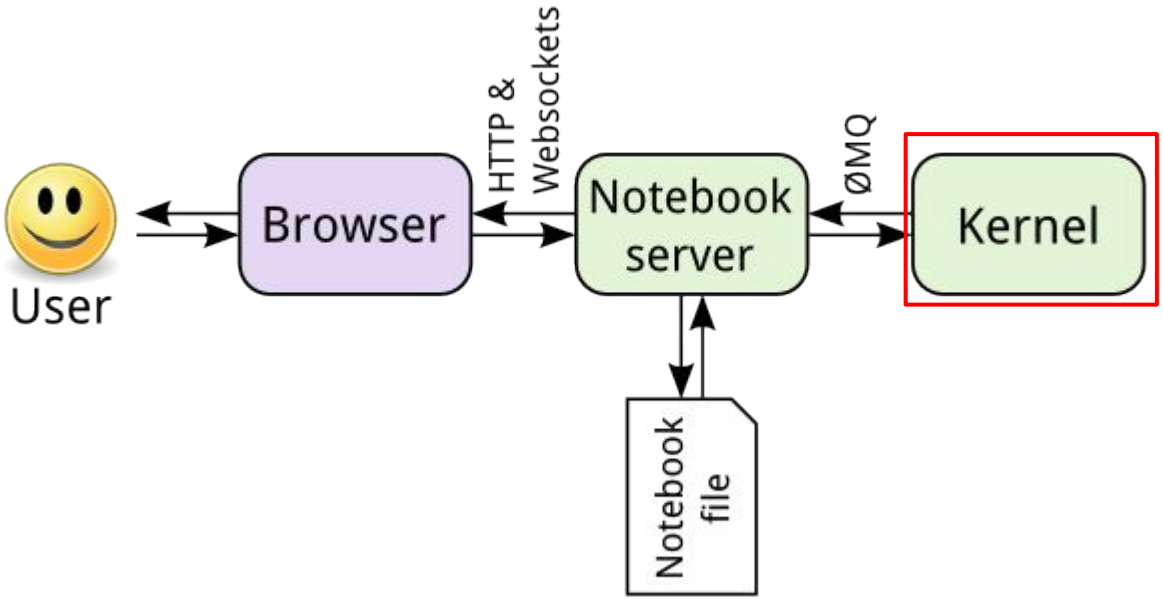
User
Browser
HTTP &
Websockets
Notebook
server
ØMQ
Kernel
Notebook
file

Kernel

- Where user code execution happens
- Separate process from the Notebook Server
- Also handles other types of requests, like code completion

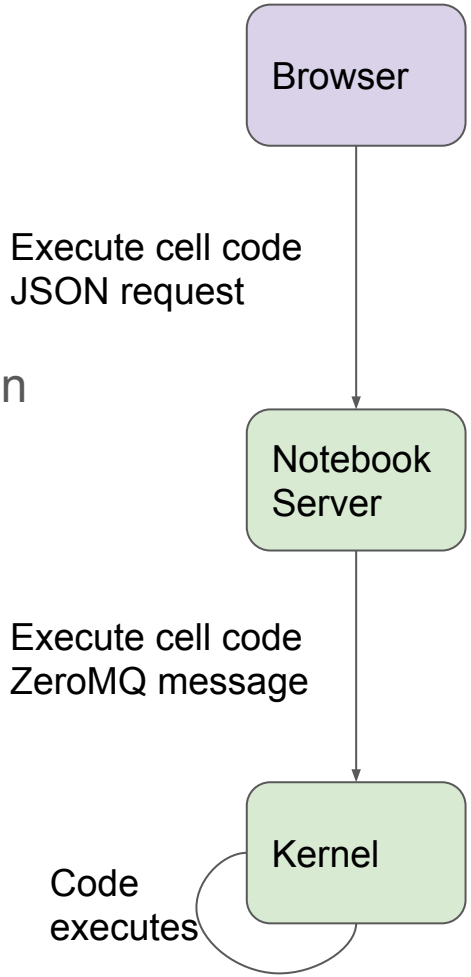

More about Kernels

- Jupyter messaging protocol
- Multiple clients can connect with different views
  - Qt console
  - Interactive shell/Terminal
  - Notebook

Support for programming languages

Kernel spec