

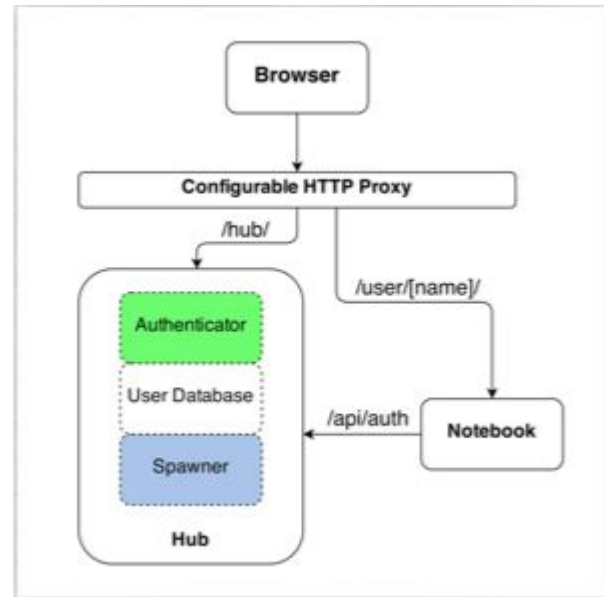
# Jupyterhub

# Jupyterhub: Jupyter as a Service

- Service to deploy notebooks in a multi-user environment
- Manages user authentication, notebook deployment and web proxies



A login form with an orange "Sign in" button at the top. Below the button are two yellow input fields: one for "Username:" containing the text "shreyas" and another for "Password:" with masked characters. At the bottom of the form is an orange "Sign in" button.



# Motivation

- ✗ Users running their own web servers on a shared cluster makes security folks very nervous.
- ✗ Difficult to support and manage different kernels and environments

## Jupyterhub to rescue




- ✓ Deploy notebooks in a standard SECURE manner
- ✓ Package known kernels out of the box
- ✓ Access shared resources on the platform
  - Filesystems, Batch Queue, Network, DBs

Thanks to Carol Willing from the Jupyter Project for the next few slides (taken from her “Thing Explainer” talk)

<https://www.slideshare.net/willingc/jupyterhub-a-thing-explainer-overview>



A WAY TO GIVE A  
JUPYTER  
NOTEBOOK SERVER  
TO EACH PERSON  
IN A GROUP OF



# WHAT IS A NOTEBOOK?


- Document
- Environment
- Web app

We have already computed  $P(X|A)$  above. On the other hand,  $P(X| \sim A)$  is subjective: our code can pass tests but still have a bug in it, though the probability there is a bug present is reduced. Note this is dependent on the number of tests performed, the degree of complication in the tests, etc. Let's be conservative and assign  $P(X| \sim A) = 0.5$ . Then

$$P(A|X) = \frac{1 \cdot p}{1 \cdot p + 0.5(1 - p)}$$
$$= \frac{2p}{1 + p}$$

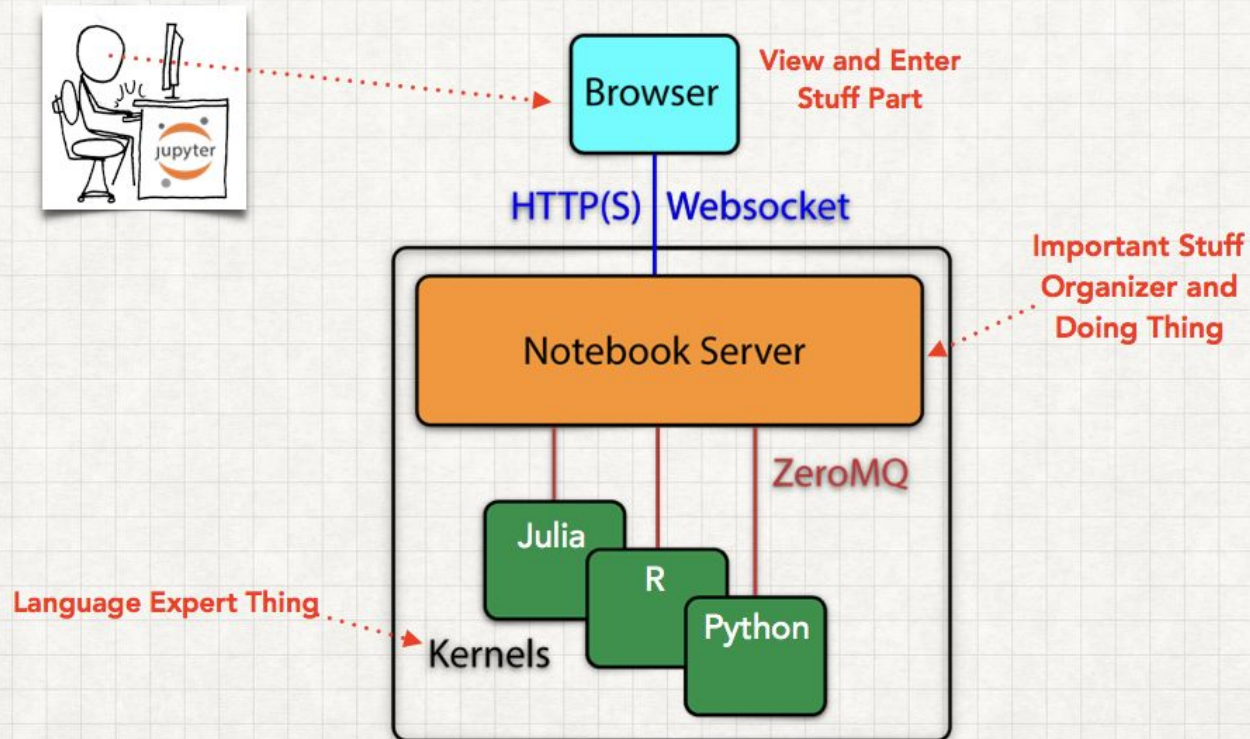
This is the posterior probability. What does it look like as a function of our prior,  $p \in [0, 1]$ ?

```
figsize(12.5, 4)
p = np.linspace(0, 1, 50)
plt.plot(p, 2 * p / (1 + p), color="#348AB0", lw=3)
# plt.fill_between(p, 2*p/(1+p), alpha=.5, facecolor="#A6D628")
plt.scatter(0.2, 2 * (0.2) / 1.2, s=140, c="#348AB0")
plt.xlim(0, 1)
plt.ylim(0, 1)
plt.xlabel("Prior, P(A) = p")
plt.ylabel("Posterior, P(A|X), with P(A) = p")
plt.title("Are there bugs in my code?")
<matplotlib.text.Text at 0x1851de650>
```



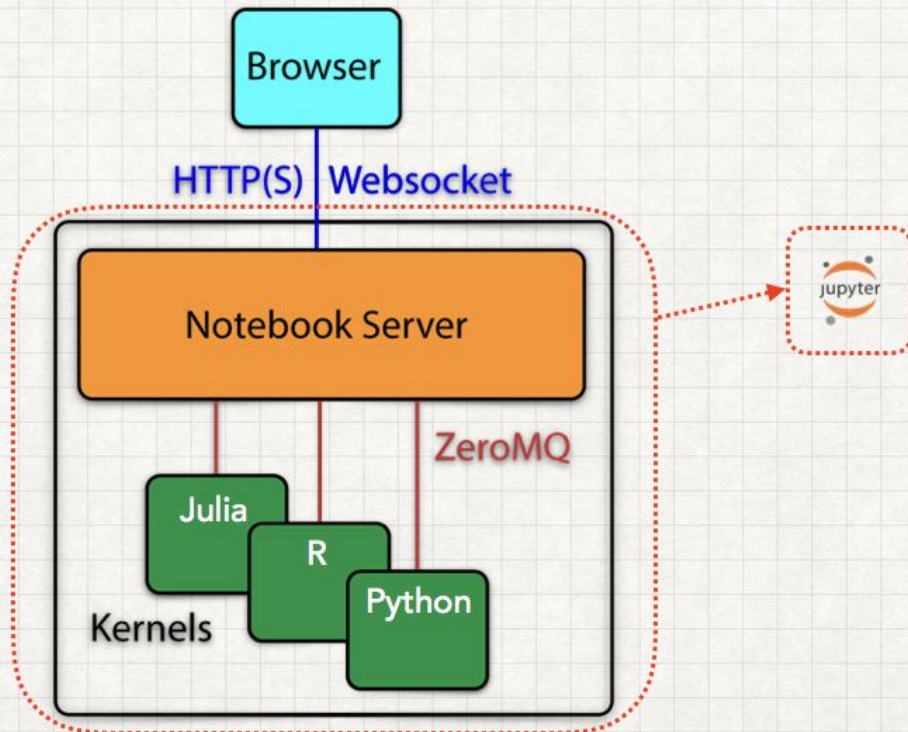
<https://github.com/CamDavidsonPilon/Probabilistic-Programming-and-Bayesian-Methods-for-Hackers>

# JUPYTER NOTEBOOK






# A SINGLE USER JUPYTER NOTEBOOK SERVER





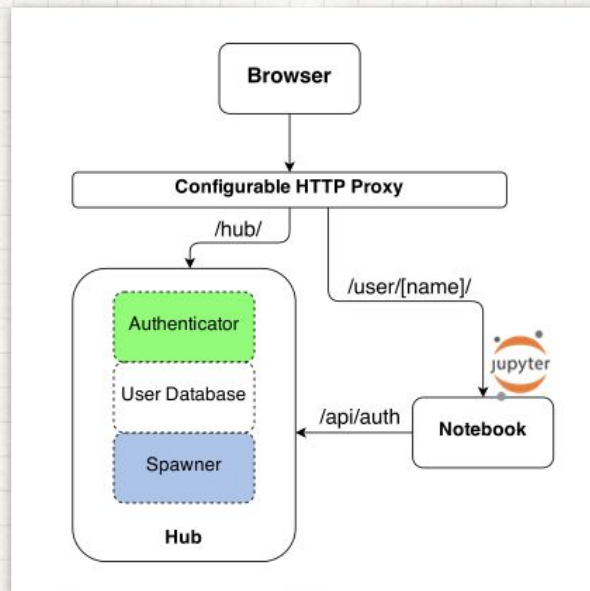


A THING TO GIVE A  
**JUPYTER**  
**NOTEBOOK SERVER**  
TO EACH PERSON  
IN A GROUP OF



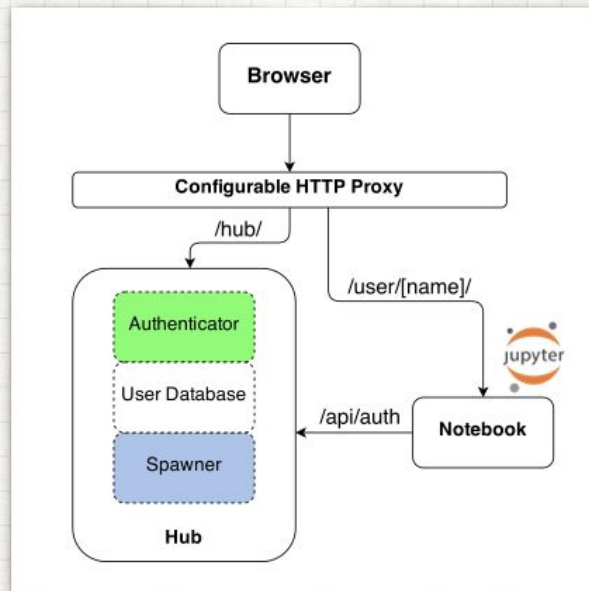
# WHAT DOES THE HUB DO?

- Manages authentication
- Spawns single-user notebook servers on-demand
- Gives each user a complete notebook server



# THE PARTS OF JUPYTERHUB

- Hub (User Database, Authenticator, Spawner)
- Users and their individual notebook servers
- Configurable HTTP Proxy



# Let's see it in action

Simple demo using docker

If you don't have docker - just follow along

This is purely a proof of concept - NOT A PRODUCTION DEMO

See link for a more in depth tutorial in a real environment

# Run the Jupyterhub Container

```
docker run -d --name jupyterhub \  
-p 8000:8000 jupyterhub/jupyterhub jupyterhub
```

# Hop Into the container

# We need to set up a couple of things ...

```
docker exec -it jupyterhub /bin/bash
```



# Install Jupyter

# Give the hub something to run on login

```
pip install jupyter
```

# Add User

```
adduser myuser
```

# In Depth Tutorial

See Min Ragan-Kelly's excellent tutorial:

<https://github.com/jupyterhub/jupyterhub-tutorial>

Includes

- Dockerspawner
- OAuth
- SSL
- kernels